

Assignment 2: Convolution

Curtis Meister

Contents

```
library(tidyverse)
library(fs)
library(keras)
library(tensorflow)

set.seed(123)
stopifnot(is_keras_available())

base_images_dir <- path_expand("~/Desktop/PetImages")
work_dir <- path_expand("~/Desktop/assignment2_runs")
dir_create(work_dir)

stopifnot(dir_exists(path(base_images_dir, "Cat")))
stopifnot(dir_exists(path(base_images_dir, "Dog")))

library(reticulate)
pil <- import("PIL.Image", delay_load = TRUE)

is_readable_image <- function(fp) {
  ok <- TRUE
  tryCatch({
    img <- pil$open(fp)
    img$load()
  }, error = function(e) {
    ok <- FALSE
  })
  ok
}

remove_corrupt_images <- function(class_dir) {
  files <- list.files(class_dir, pattern = "(?i)\\.jpg$", full.names = TRUE)
  bad <- c()
  for (f in files) {
    if (!is_readable_image(f)) bad <- c(bad, f)
  }
  if (length(bad) > 0) {
    file.remove(bad)
  }
  tibble(dir = class_dir, scanned = length(files), removed = length(bad))
}
```

```
cat_dir <- path_expand("~/Desktop/PetImages/Cat")
dog_dir <- path_expand("~/Desktop/PetImages/Dog")
```

```
clean_summary <- bind_rows(
  remove_corrupt_images(cat_dir),
  remove_corrupt_images(dog_dir)
)
```

```
clean_summary
```

```
list_class_images <- function(class_dir) {
  files <- dir_ls(class_dir, regexp = "(?i)\\.jpg$", recurse = FALSE, type = "file")
  as.character(files)
}
```

```
create_split_tree <- function(split_root) {
  dir_create(path(split_root, "train", "cats"), recurse = TRUE)
  dir_create(path(split_root, "train", "dogs"), recurse = TRUE)
  dir_create(path(split_root, "validation", "cats"), recurse = TRUE)
  dir_create(path(split_root, "validation", "dogs"), recurse = TRUE)
  dir_create(path(split_root, "test", "cats"), recurse = TRUE)
  dir_create(path(split_root, "test", "dogs"), recurse = TRUE)
}
```

```
copy_into_split <- function(files, dest_dir) {
  dir_create(dest_dir, recurse = TRUE)
  file_copy(files, dest_dir, overwrite = TRUE)
}
```

```
make_splits <- function(train_n, val_n = 500, test_n = 500, seed = 123, split_name) {
  set.seed(seed)
```

```
  cats <- list_class_images(path(base_images_dir, "Cat"))
  dogs <- list_class_images(path(base_images_dir, "Dog"))
```

```
  cats <- sample(cats)
  dogs <- sample(dogs)
```

```
  needed_per_class <- train_n + val_n + test_n
  stopifnot(length(cats) >= needed_per_class, length(dogs) >= needed_per_class)
```

```
  cats_train <- cats[1:train_n]
  cats_val <- cats[(train_n + 1):(train_n + val_n)]
  cats_test <- cats[(train_n + val_n + 1):(train_n + val_n + test_n)]
```

```
  dogs_train <- dogs[1:train_n]
  dogs_val <- dogs[(train_n + 1):(train_n + val_n)]
  dogs_test <- dogs[(train_n + val_n + 1):(train_n + val_n + test_n)]
```

```
  split_root <- path(work_dir, split_name)
  if (dir_exists(split_root)) dir_delete(split_root)
  create_split_tree(split_root)
```

```

copy_into_split(cats_train, path(split_root, "train", "cats"))
copy_into_split(dogs_train, path(split_root, "train", "dogs"))
copy_into_split(cats_val, path(split_root, "validation", "cats"))
copy_into_split(dogs_val, path(split_root, "validation", "dogs"))
copy_into_split(cats_test, path(split_root, "test", "cats"))
copy_into_split(dogs_test, path(split_root, "test", "dogs"))

split_root
}

```

```

img_size <- c(150, 150)
batch_size <- 20

make_generators <- function(split_root, augment = TRUE) {
  if (augment) {
    train_datagen <- image_data_generator(
      rescale = 1/255,
      rotation_range = 40,
      width_shift_range = 0.2,
      height_shift_range = 0.2,
      shear_range = 0.2,
      zoom_range = 0.2,
      horizontal_flip = TRUE,
      fill_mode = "nearest"
    )
  } else {
    train_datagen <- image_data_generator(rescale = 1/255)
  }

  test_datagen <- image_data_generator(rescale = 1/255)

  train_gen <- flow_images_from_directory(
    path(split_root, "train"),
    train_datagen,
    target_size = img_size,
    batch_size = batch_size,
    class_mode = "binary",
    classes = c("cats", "dogs")
  )

  val_gen <- flow_images_from_directory(
    path(split_root, "validation"),
    test_datagen,
    target_size = img_size,
    batch_size = batch_size,
    class_mode = "binary",
    classes = c("cats", "dogs")
  )

  test_gen <- flow_images_from_directory(
    path(split_root, "test"),
    test_datagen,
    target_size = img_size,

```

```

    batch_size = batch_size,
    class_mode = "binary",
    shuffle = FALSE,
    classes = c("cats", "dogs")
  )

  list(train = train_gen, val = val_gen, test = test_gen)
}

```

```

build_scratch_model <- function() {
  model <- keras_model_sequential() %>%
    layer_conv_2d(filters = 32, kernel_size = c(3,3), activation = "relu",
                  input_shape = c(img_size[1], img_size[2], 3)) %>%
    layer_max_pooling_2d(pool_size = c(2,2)) %>%

    layer_conv_2d(filters = 64, kernel_size = c(3,3), activation = "relu") %>%
    layer_max_pooling_2d(pool_size = c(2,2)) %>%

    layer_conv_2d(filters = 128, kernel_size = c(3,3), activation = "relu") %>%
    layer_max_pooling_2d(pool_size = c(2,2)) %>%

    layer_conv_2d(filters = 128, kernel_size = c(3,3), activation = "relu") %>%
    layer_max_pooling_2d(pool_size = c(2,2)) %>%

    layer_flatten() %>%
    layer_dense(units = 512, activation = "relu") %>%
    layer_dropout(rate = 0.5) %>%
    layer_dense(units = 1, activation = "sigmoid")

  model %>% compile(
    loss = "binary_crossentropy",
    optimizer = optimizer_rmsprop(learning_rate = 1e-4),
    metrics = c("accuracy")
  )

  model
}

```

```

build_vgg16_feature_extractor <- function() {
  conv_base <- application_vgg16(
    weights = "imagenet",
    include_top = FALSE,
    input_shape = c(img_size[1], img_size[2], 3)
  )

  freeze_weights(conv_base)

  model <- keras_model_sequential() %>%
    conv_base %>%
    layer_flatten() %>%
    layer_dense(256, activation = "relu") %>%
    layer_dropout(0.5) %>%
    layer_dense(1, activation = "sigmoid")
}

```

```

model %>% compile(
  loss = "binary_crossentropy",
  optimizer = optimizer_rmsprop(learning_rate = 2e-4),
  metrics = c("accuracy")
)

model
}

run_experiment <- function(approach = c("scratch", "pretrained"),
                           train_n,
                           epochs = 20,
                           seed = 123) {
  approach <- match.arg(approach)

  split_name <- paste0(approach, "_train", train_n)
  split_root <- make_splits(train_n = train_n, val_n = 500, test_n = 500,
                           seed = seed, split_name = split_name)

  gens <- make_generators(split_root, augment = TRUE)

  model <- if (approach == "scratch") build_scratch_model() else build_vgg16_feature_extractor()

  steps_per_epoch <- as.integer((2 * train_n) / batch_size)
  val_steps <- as.integer((2 * 500) / batch_size)
  test_steps <- as.integer((2 * 500) / batch_size)

  callbacks <- list(
    callback_early_stopping(
      monitor = "val_accuracy",
      patience = 3,
      restore_best_weights = TRUE
    )
  )

  history <- model %>% fit(
    gens$train,
    steps_per_epoch = steps_per_epoch,
    epochs = epochs,
    validation_data = gens$val,
    validation_steps = val_steps,
    callbacks = callbacks,
    verbose = 0
  )

  val_metrics <- model %>% evaluate(gens$val, steps = val_steps, verbose = 0)
  test_metrics <- model %>% evaluate(gens$test, steps = test_steps, verbose = 0)

  tibble(
    approach = approach,
    train_n = train_n,
    epochs = epochs,
    val_loss = as.numeric(val_metrics["loss"]),

```

```

    val_acc = as.numeric(val_metrics["accuracy"]),
    test_loss = as.numeric(test_metrics["loss"]),
    test_acc = as.numeric(test_metrics["accuracy"])
  )
}

```

```

scratch_sizes <- c(1000, 1500, 2000)
pretrained_sizes <- c(1000, 500, 1500)

results <- tibble()

for (n in scratch_sizes) {
  message("Scratch model, train_n = ", n)
  one <- run_experiment("scratch", train_n = n, epochs = 12, seed = 123)
  results <- bind_rows(results, one)
}

```

```

## Found 2000 images belonging to 2 classes.
## Found 1000 images belonging to 2 classes.
## Found 1000 images belonging to 2 classes.

```

```

## Found 3000 images belonging to 2 classes.
## Found 1000 images belonging to 2 classes.
## Found 1000 images belonging to 2 classes.

```

```

## Found 4000 images belonging to 2 classes.
## Found 1000 images belonging to 2 classes.
## Found 1000 images belonging to 2 classes.

```

```

for (n in pretrained_sizes) {
  message("Pretrained model, train_n = ", n)
  one <- run_experiment("pretrained", train_n = n, epochs = 6, seed = 123)
  results <- bind_rows(results, one)
}

```

```

## Found 2000 images belonging to 2 classes.
## Found 1000 images belonging to 2 classes.
## Found 1000 images belonging to 2 classes.

```

```

## Found 1000 images belonging to 2 classes.
## Found 1000 images belonging to 2 classes.
## Found 1000 images belonging to 2 classes.

```

```

## Found 3000 images belonging to 2 classes.
## Found 1000 images belonging to 2 classes.
## Found 1000 images belonging to 2 classes.

```

```

results

```

```

## # A tibble: 6 x 7
##   approach  train_n epochs val_loss val_acc test_loss test_acc

```

```
##   <chr>      <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 scratch    1000    12    0.597   0.671   0.587   0.682
## 2 scratch    1500    12    0.541   0.742   0.560   0.716
## 3 scratch    2000    12    0.553   0.720   0.558   0.721
## 4 pretrained 1000     6    0.261   0.885   0.240   0.905
## 5 pretrained  500     6    0.273   0.886   0.274   0.886
## 6 pretrained 1500     6    0.239   0.904   0.258   0.892
```

```
best_by_approach <- results %>%
  group_by(approach) %>%
  slice_max(val_acc, n = 1, with_ties = FALSE)
```

```
best_by_approach
```

```
## # A tibble: 2 x 7
## # Groups:   approach [2]
##   approach  train_n epochs val_loss val_acc test_loss test_acc
##   <chr>      <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 pretrained 1500     6    0.239   0.904   0.258   0.892
## 2 scratch    1500    12    0.541   0.742   0.560   0.716
```

```
results %>%
  arrange(approach, train_n) %>%
  mutate(val_acc = round(val_acc, 4), test_acc = round(test_acc, 4)) %>%
  knitr::kable()
```

approach	train_n	epochs	val_loss	val_acc	test_loss	test_acc
pretrained	500	6	0.2727180	0.886	0.2736089	0.886
pretrained	1000	6	0.2610338	0.885	0.2396473	0.905
pretrained	1500	6	0.2385759	0.904	0.2576477	0.892
scratch	1000	12	0.5973449	0.671	0.5870206	0.682
scratch	1500	12	0.5405011	0.742	0.5597976	0.716
scratch	2000	12	0.5526725	0.720	0.5578179	0.721

```
ggplot(results, aes(x = train_n, y = test_acc, group = approach)) +
  geom_line() +
  geom_point() +
  labs(
    title = "Test Accuracy vs Training Sample Size",
    x = "Training samples per class",
    y = "Test accuracy"
  )
```

